

What we check, and how each verdict is decided

The complete set of factors assessed in a GEO technical readiness scan, with the explicit rule behind every result.

35

FACTORS CHECKED

3

LAYERS

31

CONTRIBUTE TO THE SCORE

4

CONTEXT ONLY, NEVER SCORED

Why the rules are published

Generative engine optimisation is one machine reading another. A verdict about a page should therefore be reproducible by anyone holding the same page – so every check below states the exact rule that decides it.

Three things follow. Two scans of an unchanged site return the same result, which is what makes progress between audits measurable rather than anecdotal. Nothing rests on the judgement of whoever happened to run the scan. And where a factor genuinely cannot be settled remotely it is reported as **unverified** and excluded from the score entirely.

This document is generated directly from the scanner own check registry. It cannot describe a check the tool does not run, or omit one it does.

Layer 1 — Access

CHECKS
11

Can AI systems reach this site at all?

SCORING
reported separately

Whether crawlers and retrieval bots can fetch the site, and whether anything blocks, redirects or rate-limits them. Nothing downstream matters if this fails.

Who can act on it: Hosting, DNS and CDN – usually the host or platform, not the web team.

1	Site reachability from the scanner	site-reachability	NOT SCORED
WHAT IS MEASURED	Whether any request to the origin completes, and the transport-level reason when none do.		
RULE	Reported as UNVERIFIED whenever every user-agent fails at the connection layer. Never a failure of the site: a refused connection is most often edge protection rejecting the scanner's network, and the site typically serves real visitors normally throughout.		
WHY IT MATTERS FOR CITATION	Distinguishes "we could not measure this" from "this site is broken". Scoring an unreachable site would publish our network conditions as a verdict about the client.		

2 **www / non-www resolution**

host-variant

WHAT IS MEASURED	Whether both the apex domain and the www hostname resolve and serve the site.
RULE	FAIL when one form resolves and the other does not. The scan continues against whichever answers.
WHY IT MATTERS FOR CITATION	Anyone typing or linking the missing form reaches nothing, and engines see two hostnames rather than one consolidated site.

3 **Multi-user-agent crawl test**

multi-ua

WHAT IS MEASURED	The homepage fetched with seven user-agents: a browser, GPTBot, ClaudeBot, Googlebot, OAI-SearchBot, PerplexityBot and a plain default.
RULE	FAIL only when the server returns an explicit refusal (401, 402, 403, 405, 451) to a bot while the browser request succeeds. Rate limiting (429) and dropped connections WARN instead — they are transient, and scanning can cause them.
WHY IT MATTERS FOR CITATION	Catches user-agent-based bot blocking at the host. Requires a deliberate refusal status: a timeout is evidence of an unreliable connection, not of a policy.

4 **robots.txt AI crawler rules**

robots-txt

WHAT IS MEASURED	18 crawler tokens evaluated against robots.txt, split into live retrieval, model training and search indexing.
RULE	FAIL if any retrieval or search crawler is blocked (OAI-SearchBot, ChatGPT-User, Claude-User, Claude-SearchBot, PerplexityBot, Perplexity-User, Google-Extended, Googlebot, Bingbot, *). WARNING if only training crawlers are blocked (GPTBot, ClaudeBot, CCBot, Applebot-Extended and similar).
WHY IT MATTERS FOR CITATION	Blocking retrieval means the site cannot be cited in answers at all. Blocking training is a legitimate licensing choice with a slower cost, so it is named rather than treated as a defect. Google-Extended is singled out because the one-click AI toggle in Yoast and RankMath sets it, removing sites from Gemini without their owners knowing.

5 **X-Robots-Tag header**

x-robots-tag

WHAT IS MEASURED	The X-Robots-Tag HTTP response header.
RULE	FAIL if it contains noindex or none. WARNING if present without those directives.
WHY IT MATTERS FOR CITATION	A server-level indexing block set in .htaccess, nginx config or a plugin — invisible in the page source, and easily missed by hand.

6 **Meta robots noindex tag**

noindex-meta

WHAT IS MEASURED	meta name="robots" and meta name="googlebot" in the document head.
RULE	FAIL if either contains noindex.
WHY IT MATTERS FOR CITATION	Commonly left on after a staging launch, or by the WordPress "discourage search engines" setting.

7 XML sitemap

sitemap

WHAT IS MEASURED

Sitemap declarations in robots.txt, then five common paths: /sitemap.xml, /sitemap_index.xml, /wp-sitemap.xml, /sitemap-index.xml, /sitemap1.xml.

RULE

PASS on a valid sitemap with 3 or more URLs. WARNING if none is found, or fewer than 3 URLs are listed. FAIL if a sitemap exists but is empty or malformed.

WHY IT MATTERS FOR CITATION

Speeds discovery for every crawler. Checking declarations first avoids reporting a correctly configured site as having no sitemap.

8 Response time

response-time

WHAT IS MEASURED

The fastest of two isolated homepage requests, each issued with nothing else contending for the origin.

RULE

WARNING above 2500ms. Never a FAIL – a timing sample should not be able to gate anything. UNVERIFIED if no request completed.

WHY IT MATTERS FOR CITATION

Slow origins get abandoned by crawlers. Two isolated samples rather than one keeps the figure stable enough to compare across re-audits.

9 Edge bot protection (CDN/WAF)

edge-protection

NOT SCORED

WHAT IS MEASURED

CDN and WAF fingerprints in response headers – Cloudflare, Sucuri, Akamai, Fastly, Imperva.

RULE

Always UNVERIFIED when a provider is detected. Edge services identify verified bots by source IP range, not user-agent string, so user-agent testing cannot settle whether AI crawlers are permitted – it passes regardless.

WHY IT MATTERS FOR CITATION

The scanner's clearest blind spot, reported as such rather than hidden behind a pass. Carries the exact manual check, and names the hosting platform when the CDN belongs to the platform rather than the client.

10 Hosting platform

hosting-platform

NOT SCORED

WHAT IS MEASURED

Platform fingerprints in the page and headers – Lovable, Wix, Squarespace, Framr, Webflow, Shopify, self-hosted WordPress.

RULE

Informational. Never a pass or a failure.

WHY IT MATTERS FOR CITATION

Decides who can fix what. On a managed platform the CDN belongs to the host, so advice to change bot settings sends the client looking for a dashboard they do not have.

11 llms.txt

llms-txt

NOT SCORED

WHAT IS MEASURED

Presence of /llms.txt.

RULE

Presence only. Its absence is never counted against a site.

WHY IT MATTERS FOR CITATION

An emerging convention that no major engine has committed to honouring. Reported as an early-mover signal, not a requirement.

Layer 2 — Machine readability

CHECKS

15

Can a machine parse the page and understand its structure?

SCORING

counts toward score

Whether the content exists in the HTML without JavaScript, and whether titles, metadata, structured data, headings and landmarks describe it in a form a machine can interpret.

Who can act on it: Developer or CMS.

1 Content visible without JavaScript

js-rendering

WHAT IS MEASURED

Readable words in the server response, script count, and whether a framework mount point arrived empty.

RULE

FAIL if the server returns under 60 words alongside scripts, or a mount point (`#root`, `#__next`, `#app`) contains under 30 words.

WHY IT MATTERS FOR CITATION

Most AI crawlers do not execute JavaScript. If this fails, every other on-page result for the page is unreliable — the checks are reading an empty shell, not the real page.

2 Title tag

title

WHAT IS MEASURED

Presence and length of the title element.

RULE

FAIL if missing. WARNING above 60 characters.

WHY IT MATTERS FOR CITATION

The strongest single statement of what a page is about, used by both search and AI surfaces.

3 Meta description

meta-description

WHAT IS MEASURED

Presence and length of the meta description.

RULE

FAIL if missing. WARNING above 160 characters.

WHY IT MATTERS FOR CITATION

Frequently lifted verbatim as the summary shown alongside a citation.

4 Schema.org / JSON-LD

schema

WHAT IS MEASURED

JSON-LD blocks parsed for @type, compared against eight common types.

RULE

FAIL if no structured data is present. WARNING if more than four common types are absent.

WHY IT MATTERS FOR CITATION

The most direct way to state machine-readable facts about an entity. Among the highest-leverage signals available for citation.

5 Heading structure

headings

WHAT IS MEASURED

Count of h1 elements and whether heading levels are skipped.

RULE

PASS with exactly one h1 and no skipped levels. FAIL with no h1. WARNING otherwise.

WHY IT MATTERS FOR CITATION

Headings are how a machine segments a page into passages it can retrieve individually.

6 Sequential heading hierarchy

heading-hierarchy

WHAT IS MEASURED	Whether heading levels descend without gaps.
RULE	WARNING when a level is skipped (h1 straight to h3).
WHY IT MATTERS FOR CITATION	A broken hierarchy produces a misleading document outline, blurring where one topic ends and the next begins.

7 Canonical tag

canonical

WHAT IS MEASURED	link rel="canonical" and whether it is self-referencing or cross-domain.
RULE	FAIL if it points to another domain. WARNING if absent.
WHY IT MATTERS FOR CITATION	A canonical left pointing at staging after a migration can suppress the live page entirely.

8 Open Graph tags

open-graph

WHAT IS MEASURED	og:title, og:description and og:type.
RULE	FAIL if all three are missing. WARNING if some are.
WHY IT MATTERS FOR CITATION	Controls how the page is represented when shared or previewed, including by tools that fetch preview metadata.

9 Image alt text coverage

image-alt

WHAT IS MEASURED	Proportion of img elements carrying non-empty alt text.
RULE	PASS at 80% or above. WARNING from 40%. FAIL below 40%.
WHY IT MATTERS FOR CITATION	Alt text is the only description of an image a non-visual system receives.

10 Main landmark presence

main-landmark

WHAT IS MEASURED	A main element or role="main".
RULE	FAIL if absent.
WHY IT MATTERS FOR CITATION	Tells an agent which part of the document is the content, as distinct from navigation and chrome.

11 Form input labels

form-labels

WHAT IS MEASURED	Whether form inputs have associated labels or accessible names.
RULE	WARNING when inputs lack them.
WHY IT MATTERS FOR CITATION	An AI agent completing a form on a user's behalf can only identify a field by its accessible name.

12 Accessibility tree health

a11y-tree-health

WHAT IS MEASURED	Composite of the landmark, heading hierarchy, form label and image alt results.
RULE	Derived from its component checks.
WHY IT MATTERS FOR CITATION	The accessibility tree is broadly what an agentic browser navigates by, so its health predicts how well an agent can operate the page.

13	Contact info machine-readability	contact-machine-readable
WHAT IS MEASURED	tel: and mailto: links, and ContactPoint schema. Only applied when the page shows contact details at all.	
RULE	WARNING when contact details appear as plain text only.	
WHY IT MATTERS FOR CITATION	Lets an agent act on the details rather than merely read them.	

14	Structured data completeness	schema-completeness
WHAT IS MEASURED	Every property actually present on each recognised schema.org type, checked against that type's required and recommended properties.	
RULE	FAIL when a type is missing a required property (Organization without name or url, LocalBusiness without an address). WARNING when most recommended properties are absent. Only recognised types are assessed.	
WHY IT MATTERS FOR CITATION	Presence alone proves nothing. A block carrying only @type and a name gives an engine no way to tie the markup to a real entity, yet it passes a presence check exactly as a complete one does.	

15	FAQ schema vs visible content	faq-schema-match
WHAT IS MEASURED	FAQPage schema questions matched against text visible on the page. Only applied where FAQ schema exists.	
RULE	PASS when 70% or more of schema questions appear in the visible content.	
WHY IT MATTERS FOR CITATION	Structured data that does not reflect the page can be treated as manipulative.	

Layer 3 — Substance & authority

Is there anything here worth citing, and any reason to trust it?

CHECKS

9

SCORING

counts toward score

Whether the page carries enough specific, well-scoped content to be quoted, and whether it demonstrates a real, accountable organisation behind it. A page can pass every technical check and still never be cited because there is nothing in it to quote.

Who can act on it: Content and marketing.

1	Content depth for citation	content-depth
WHAT IS MEASURED	Words of main content, with navigation, header and footer removed.	
RULE	FAIL under 300 words. WARNING under 600. PASS at 600 or above; 1200 or above is recorded as substantial.	
WHY IT MATTERS FOR CITATION	Below roughly 300 words there is not enough for an engine to extract a useful passage. A page can be crawled perfectly and never be quoted because there is nothing in it worth quoting.	

2 Page scope — should this be split?

page-scope

WHAT IS MEASURED	Content split at h2 boundaries and measured per section, plus whether navigation is built entirely from in-page anchors.
RULE	FAIL when navigation is anchor-only across 3 or more topics — a single-page site. WARNING at 4 or more sections averaging under 120 words.
WHY IT MATTERS FOR CITATION	Engines retrieve and cite at page level. Topics sharing one document cannot be returned for queries about them individually, however good the writing.

3 Authority & credibility signals

authority-signals

WHAT IS MEASURED	Six signals, detected bilingually (English and Spanish): named people, client evidence, credentials, physical address, direct contact links, external profile links including sameAs.
RULE	PASS at 5 of 6 or more. WARNING at 3 or 4. FAIL at 2 or fewer.
WHY IT MATTERS FOR CITATION	Engines favour sources visibly belonging to a real, accountable organisation. Without these a site reads as anonymous marketing copy however well it is built.

4 Answer-shaped content

answer-format

WHAT IS MEASURED	Question-form headings (English and Spanish), list items, tables and definition lists.
RULE	PASS with 2 or more question-form headings, or a structure score of 12 or more (lists count 1, tables 5, definition lists 3).
WHY IT MATTERS FOR CITATION	Retrieval matches a question to a passage that answers it. Q&A sections, lists and comparison tables are markedly easier to lift an answer from than continuous prose.

5 Named entity / specificity signal

entities

WHAT IS MEASURED	Proper-noun phrases, numbers and year references in the main content.
RULE	WARNING when fewer than 3 proper nouns and fewer than 3 numbers are present.
WHY IT MATTERS FOR CITATION	Specific, checkable facts are what gets quoted. Generic marketing language offers nothing to cite.

6 First-250-words specificity

first-250-specificity

WHAT IS MEASURED	The same entity heuristic applied to the opening 250 words only.
RULE	PASS with 3 or more proper nouns, or 3 or more numbers, in the opening.
WHY IT MATTERS FOR CITATION	Extraction weights the beginning of a page more heavily than material further down.

7 Author / credential attribution

author-attribution

WHAT IS MEASURED	Bylines, rel="author" links, and Person schema tied to authorship.
RULE	WARNING when none is found.
WHY IT MATTERS FOR CITATION	Named, credentialed sources are weighted more heavily than unattributed corporate copy.

8 Freshness signals

content-freshness

WHAT IS MEASURED	dateModified, datePublished and uploadDate in structured data, <time datetime> elements, and article date metadata.
RULE	WARNING when no machine-readable date exists at all, or when the most recent is over a year old. PASS within a year.
WHY IT MATTERS FOR CITATION	Engines weight recency. A page carrying no readable date cannot be assessed for it, which is itself a disadvantage against a competitor whose page is visibly current.

9 Boilerplate / templated content

boilerplate

WHAT IS MEASURED	Pairwise five-word shingle similarity between analysed pages.
RULE	WARNING above 60% similarity between any pair. UNVERIFIED with fewer than two pages, since there is nothing to compare.
WHY IT MATTERS FOR CITATION	Near-identical pages compete with each other and signal low-value templating.